

DecisionOS Report: decision_8792883d99

Pack: Product | Created: 2026-06-28T06:28:42.983687

Executive Summary

Launch AI workspace module should be treated as a decision system, not a one-off task. The current recommendation is: Proceed with staged execution. The RICE weighted score is 83/100, based on customer pain intensity, product differentiation, adoption friction, retention path.

Decision Verdict

Proceed with staged execution: choose 'Launch MVP' as the lead path, with evidence gates tied to the RICE score.

Scoring Engine

RICE: 83/100

- Reach: 81/100 | weight 30% | Strong supporting signals
- Impact: 81/100 | weight 30% | Strong supporting signals
- Confidence: 72/100 | weight 20% | Moderate support with validation needs
- Effort: 100/100 | weight 20% | Strong supporting signals

Reference Cases

Notion AI launch (Notion, 2023) - 90%

Lesson: AI features work best when they reduce friction in an already frequent workflow.

Slack product-market fit validation (Slack, 2013) - 86%

Lesson: Early PMF is visible through repeated team usage.

Figma collaboration expansion (Figma, 2016) - 80%

Lesson: A product can win by changing the collaboration model.

Core Value

- turn user demand into a prioritized product decision with measurable adoption signals
- Align stakeholders around the objective: Decide whether to launch now.
- Reduce avoidable execution waste by separating assumptions, constraints, and measurable proof.

Benchmark

- Scoring framework: RICE.
- Compare the leading option against pack dimensions: customer pain intensity, product differentiation, adoption friction, retention path.
- Current weighted score: 83/100.

- A strong decision should be specific, measurable, reversible where possible, and owned by a named team.

Execution Plan

- Frame the decision: confirm owner, objective, options, constraints, and non-negotiables.
- Collect evidence: interview users or stakeholders, review financial impact, and test feasibility.
- Score options: compare upside, downside, time-to-proof, cost, reversibility, and strategic fit.
- Run a controlled pilot: execute the smallest version that can prove or disprove the main thesis.
- Decide the next gate: scale, revise, pause, or stop based on agreed metrics.

Budget

- Current planning budget: 180,000.
- Reserve 60% for core execution, 20% for validation and measurement, 20% for contingency.
- Do not unlock the next budget tranche until the validation gate is passed.

Next Actions

- Create a one-page decision brief with owner, objective, options, constraints, and metrics.
- Select the lead option and one fallback option.
- Review scoring dimensions below 70 and fill the evidence gaps.
- Define the validation gate and schedule the decision review.
- Assign owners for evidence, execution, finance, and risk.

Risk Matrix

- Evidence quality: Medium/High. Define the smallest proof package before committing full resources.
- Execution dilution: Medium/Medium. Limit the first cycle to one owner, one operating metric, and one review cadence.
- Limited engineering capacity: Unknown/High. Convert open risks into named assumptions with validation dates.

Timeline

- Frame (Day 1-2): Decision brief and option set
- Validate (8 weeks): Evidence review and pilot design
- Commit (Next review cycle): Scale, revise, pause, or stop decision

This Week Action List

- [High] Monday: Complete a one-page decision brief with objective, options, constraints, and decision criteria. Owner: Product, engineering, sales. Duration: 2 hours.
- [High] Tuesday: Interview 3 key users or stakeholders to fill the highest-risk evidence gap. Owner: Evidence owner. Duration: Half day.

- [Medium] Wednesday: Score the lead path and fallback path by cost, risk, reversibility, and time-to-proof. Owner: Operations owner. Duration: 2 hours.
- [High] Thursday: Design the smallest validation pilot with pass/fail gates. Owner: Execution owner. Duration: Half day.
- [Medium] Friday: Run the decision review and choose scale, revise, pause, or stop. Owner: Final decision owner. Duration: 1 hour.